


# Cross-Validation & Overfitting



# Underfitting and Overfitting Overview

- A model that fits the training data too well (has a low training error) may have a higher generalization error than a model with a higher training error


Underfitting



Overfitting



# Example Data Set



**Two class problem:**

**+ : 5400 instances**


- 5000 instances generated from a Gaussian centered at (10,10)
- 400 noisy instances added

**o : 5400 instances**


- Generated from a uniform distribution

**10 % of the data used for training and 90% of the data used for testing**

## Learning Curve - Increasing number of nodes in Decision Trees

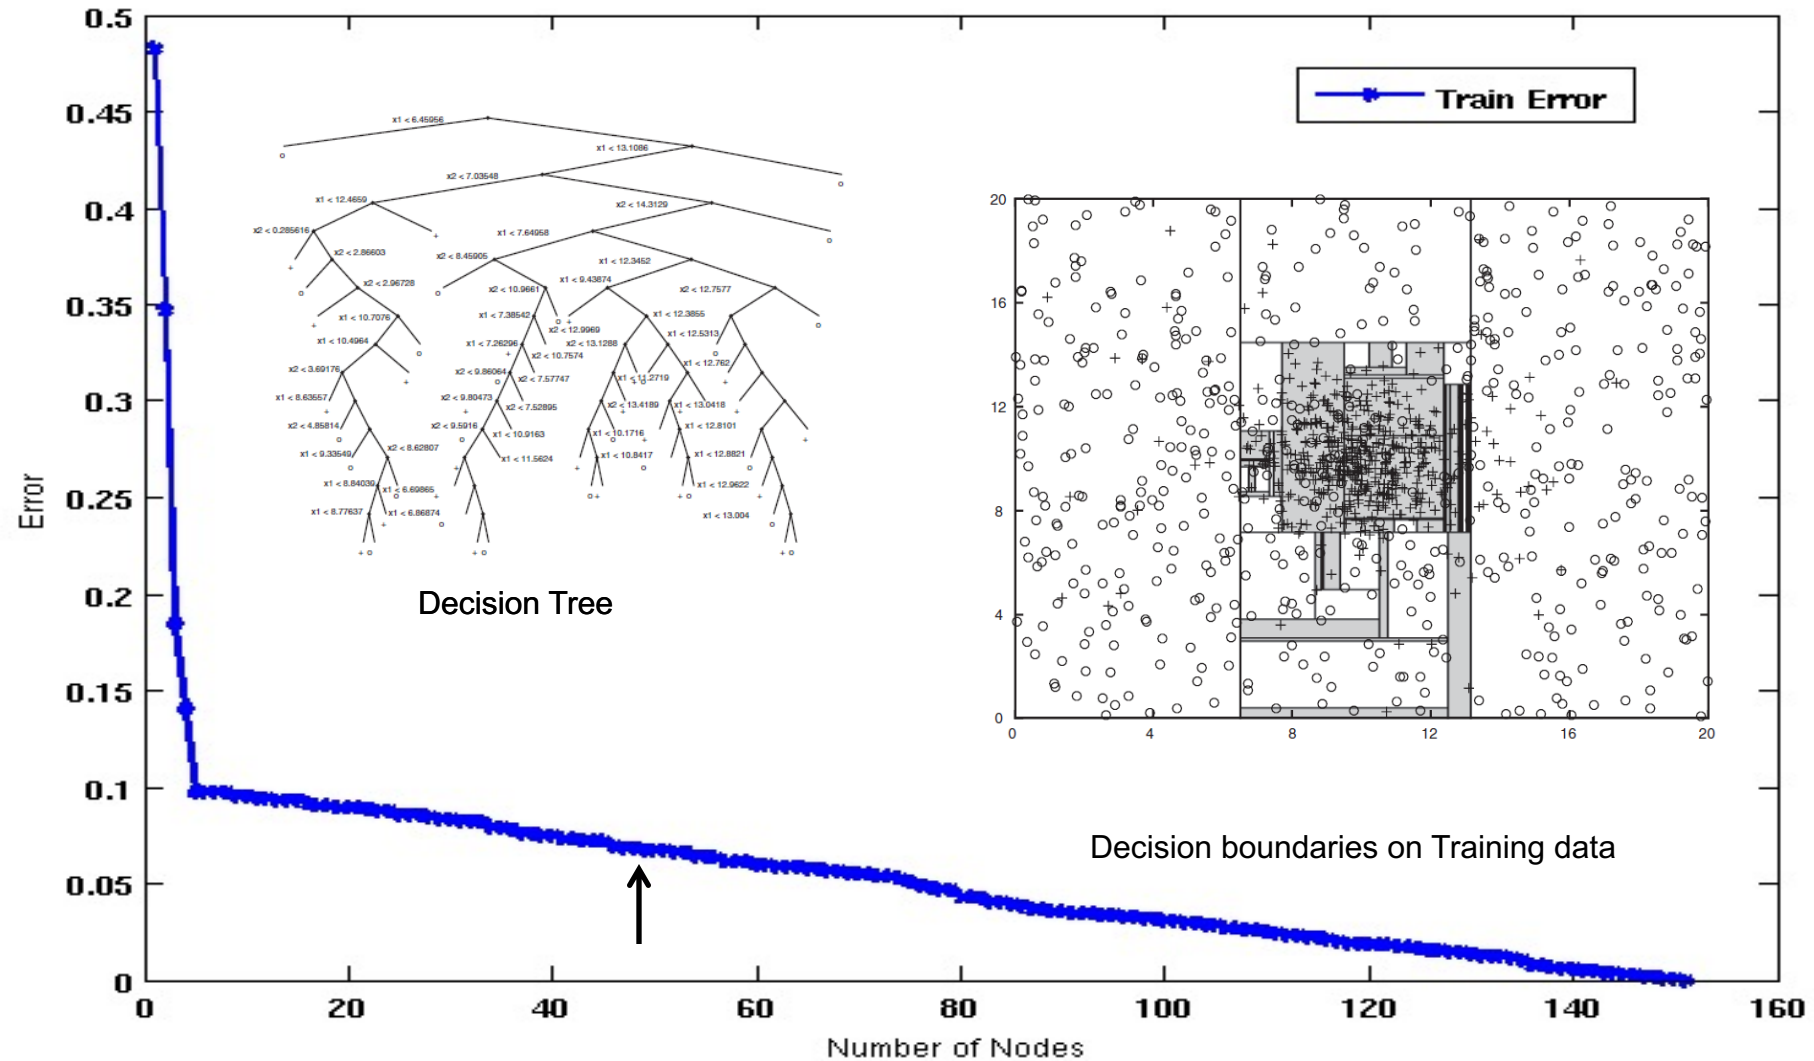


## Decision Tree with 4 nodes

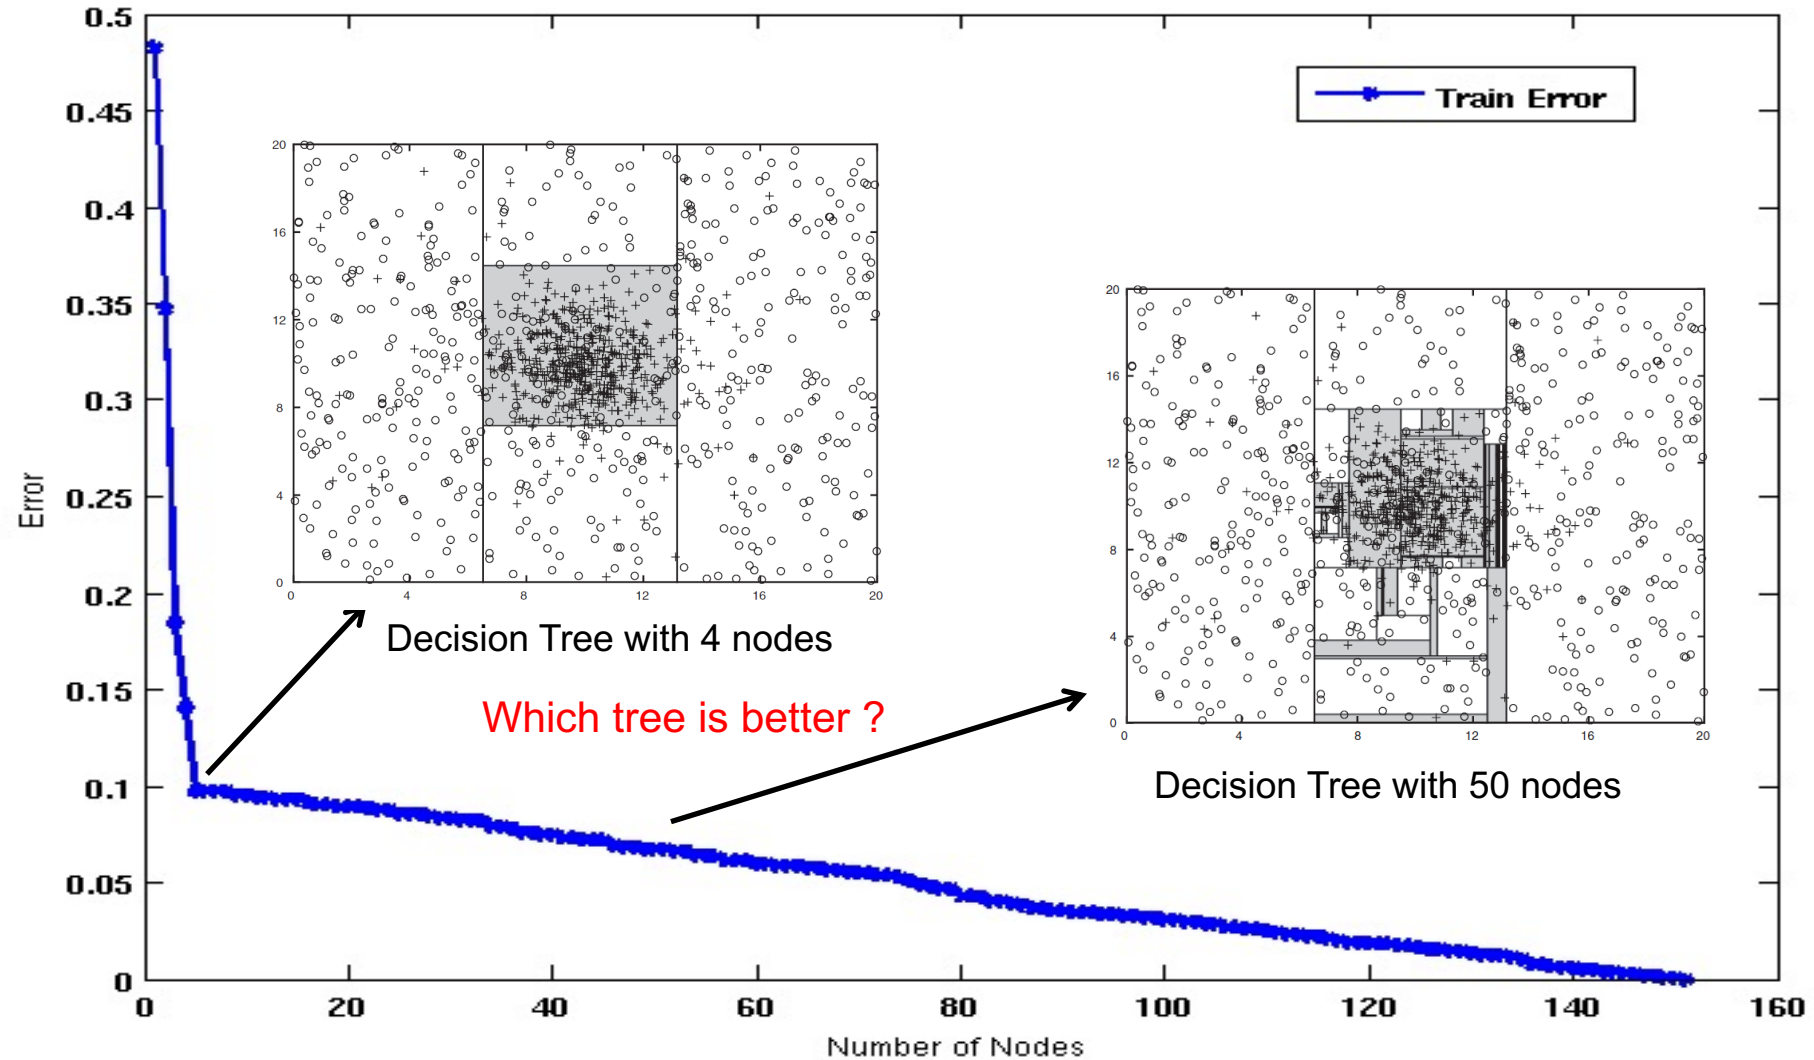




## Decision Tree with 50 nodes





Which tree is better?





# Model Overfitting





- As the model becomes more and more complex, test errors can start increasing even though training error may be decreasing

Underfitting: when model is too simple, both training and test errors are large

Overfitting: when model is too complex, training error is small but test error is large

# Model Overfitting



Using twice the number of data instances

- Increasing the size of training data reduces the difference between training and testing errors at a given size of model

# Reasons for Model Overfitting

- Limited Training Size
- High Model Complexity

# Measuring Error

How to detect under/overfitting

# Types of Error


## Training Phase

- **Training Error:** the percent of misclassification errors on the training data set

## Testing Phase

- **Test Error / Generalization Error:** the percent of misclassification errors on the test data set (previously unseen records)
- **Accuracy:** the percent of correct classifications on the test data set ( $100 - \text{Generalization Error}$ )

# Partitioning Data to Predict Error




- Holdout Method: split data into training set and test set
- Build model on training set, evaluate generalization error on test set




# Issues with Holdout Method

- Less data for training (and less data for testing)
- Some records never get trained on (and some never get tested on)
- A class overrepresented in one set will be underrepresented in the other set
- Varying performance of the model, depending on which records were held out for testing



# Issues with Holdout Method

- Less data for training (and less data for testing)
- Some records never get trained on (and some never get tested on)
- A class overrepresented in one set will be underrepresented in the other set
- Varying performance of the model, depending on which records were held out for testing




# K-Fold Cross-Validation

- Split data into K folds
  - Model building (training) and error estimation (testing) is repeated K times
  - Each iteration, one of the folds is used for testing, the rest are used for training.
  - Get an error estimate from each fold. Average them to establish a final estimate of generalization error.

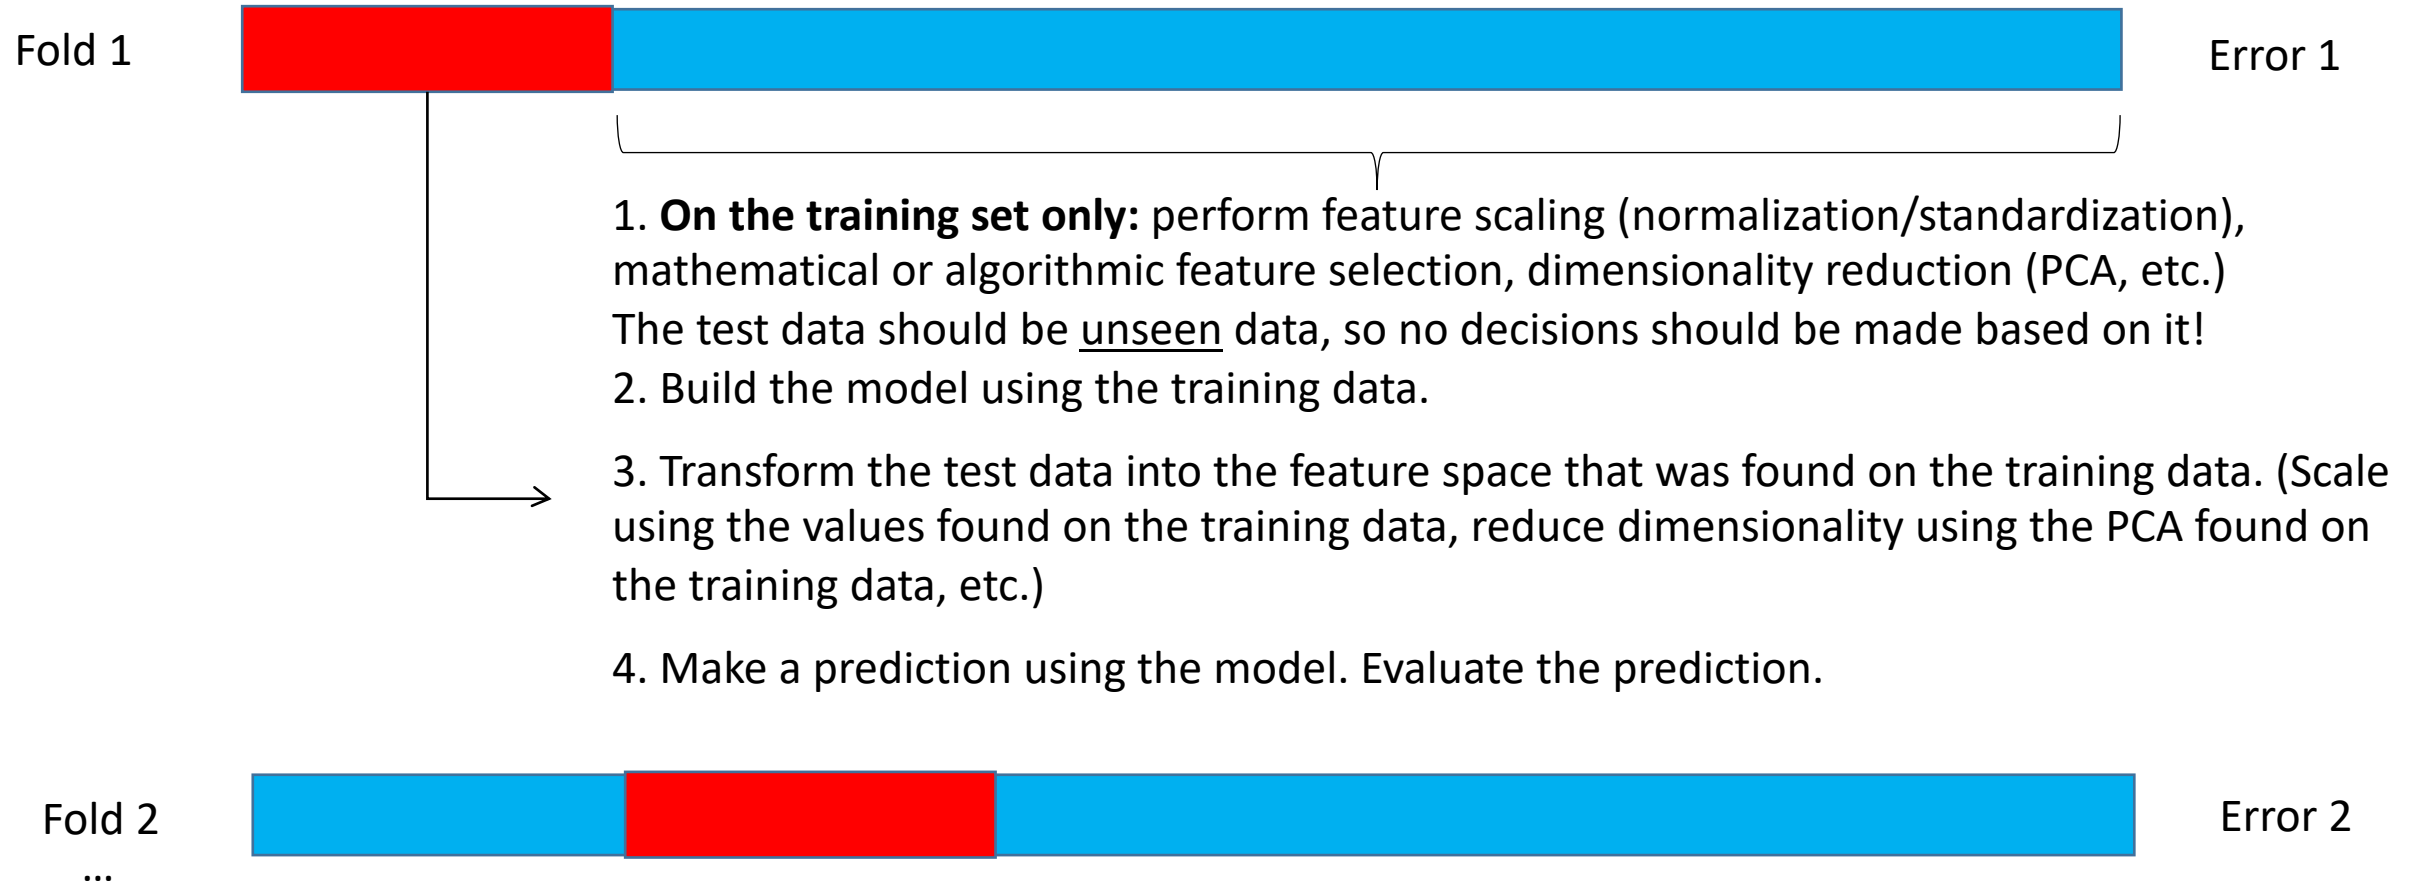


# Example: Three Fold Cross Validation



# Inside the Cross-Validation Loop

- Cross-validation is used to assess the performance of the process
- It is critical that the test data remains totally unseen and is not used for any part of the model creation. (That would be cheating! Also called: [data leakage](#).)



# Building the Final Model

- A good solution is found (you got a high accuracy) - now what...?
- The cross-validation created K different models of the data – which one do you use?
- None of them! Cross-validation was necessary to evaluate the process. Once the process is deemed good, you run the whole process on your entire dataset to make a final model.
  - Perform scaling, feature selection, dimensionality reduction, etc. on the entire dataset, and train a final model on the entire dataset (there is no test set).
- This final model will be used in the real-world to make predictions on new data.



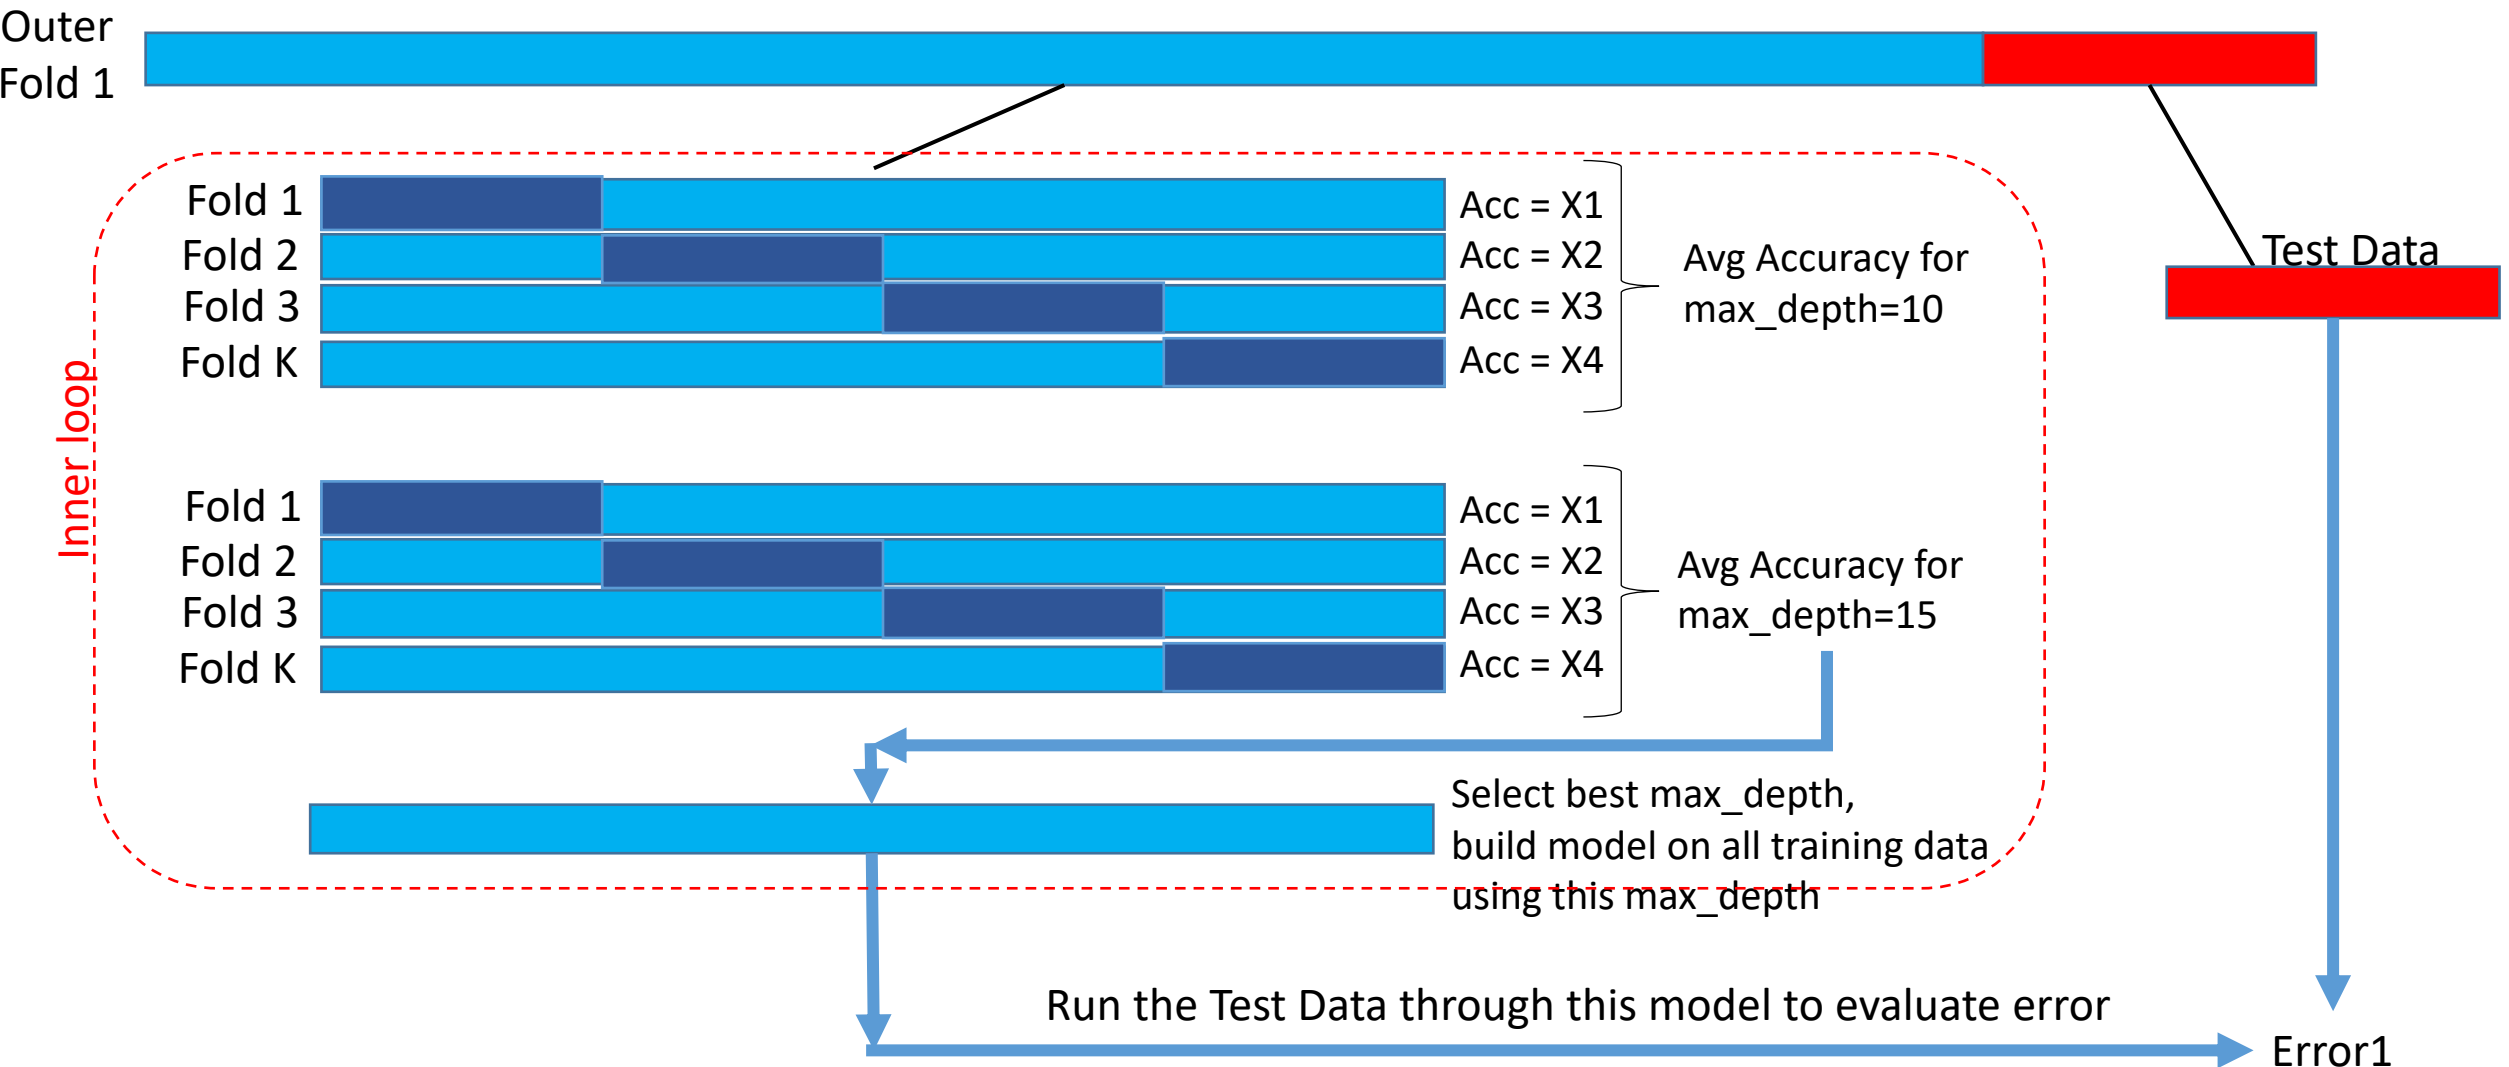
# Finding the best Hyperparameters

Nested Cross Validation

# Hyperparameters

- From the book:
- “Hyper-parameters are parameters of learning algorithms that need to be determined before learning the classification model”
- Look at the documentation for scikit-learn decision tree:
  - <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>

# Nested Cross Validation



# Choose the Hyperparameter

- For each fold of your outer loop, you calculated the best hyperparameter, so you have  $k$  votes
  - In the best-case scenario, all folds vote for the same hyperparameter value
  - In the worst-case scenario, all folds vote for a different hyperparameter and more study is required.
- Use the consensus (most votes) to select the best hyperparameter to train the final model.